

Research Plan

Decoding Diabetic Peripheral Neuropathy Heterogeneity: A Hybrid ML System Using a Patient-Personalized Vector for Disease Representation and Stratification

Student :

Mariam Mohamed Ahmed Elelidy

Supervisor :

Dr. Reda Abdelbaset Elazab Mohamed, PhD

Co-Principal Investigator at the School of Science and Engineering,
American University in Cairo (AUC), Cairo, Egypt.

Qualified Scientist:

Dr. Omnia Mugeeb Saleh Abd-Elrahman, PhD

Consultant of Endocrinology and Diabetes at the National Research
Centre (NRC), Cairo, Egypt.

Index

No.	Item	Page No.
1.	Introduction & Background	2
2.	Rationale & Significance	4
3.	Research Objectives & Expected Outcomes	6
4.	Research Questions & Hypotheses	7
5.	Methodology	8
6.	Human Participants	12
7.	Data Analysis	13
8.	Risk Assessment & Safety Considerations	16
9.	Ethical Considerations	18
10.	Translational Consideration	18
11.	Bibliography	19

1. Introduction & Background

Diabetic Peripheral Neuropathy (DPN) is among the most prevalent and disabling chronic complications of diabetes mellitus, affecting approximately **45–50%** of individuals with long-standing diabetes. Current estimates indicate that more than **130 million people worldwide** live with DPN, with prevalence expected to rise in parallel with the global diabetes epidemic. In the United States, DPN affects approximately **8–10 million adults**, while low- and middle-income countries, including Egypt, are experiencing rapidly increasing incidence driven by **delayed diagnosis** and **limited access to specialized neurological care**.

Clinically, DPN manifests as **progressive peripheral nerve injury**, leading to **neuropathic pain, sensory loss, impaired proprioception, muscle weakness, gait instability, and balance dysfunction**. These impairments significantly elevate the risk of downstream complications. Epidemiological data indicate that **15–25%** of patients with DPN develop diabetic foot ulcers during their lifetime, and that DPN accounts for **over 70%** of non-traumatic lower-limb amputations in individuals with diabetes. Beyond physical morbidity, DPN is strongly associated with **sleep disruption, depression, reduced mobility, and loss of functional independence**, resulting in substantial deterioration in quality of life.

The **economic burden** of DPN is equally substantial. Annual **direct medical costs per patient** range from approximately **USD 5,000 to 25,000**, depending on disease severity and complication profile. Globally, healthcare expenditures attributable to DPN exceed **USD 50 billion annually**, with nearly **USD 10 billion** spent each year in the United States alone. In Egypt, diabetes-related neuropathic complications contribute an estimated **USD 0.5–1 billion annually** to national healthcare spending, a burden projected to escalate as diabetes prevalence increases.

Despite its high prevalence and impact, DPN is **not a homogeneous disease**. Patients with similar diabetes duration, glycemic control, and symptom severity frequently exhibit **markedly different trajectories of nerve degeneration, symptom progression, and treatment response**. This **biological and functional heterogeneity** represents a central challenge in DPN management.

Limitations of Current DPN Stratification Approaches

Current diagnostic and staging tools—including **nerve conduction studies, symptom-based scoring systems, and clinical neurological examinations**—are primarily designed to **detect neuropathy**, not to **interpret underlying disease mechanisms**. As a result, DPN stratification remains largely **descriptive and severity-based**, typically classifying patients into broad categories such as mild, moderate, or severe neuropathy.

Key limitations include:

- Patients with **identical clinical scores** may exhibit **substantially different progression rates**
- **Symptom-based assessments** poorly predict future functional decline
- Existing tools focus on **detection rather than interpretation**
- **Biological heterogeneity** driving disease progression is largely ignored

Consequently, treatment selection commonly relies on **trial-and-error strategies**, with both pharmacological and non-pharmacological interventions adjusted reactively based on symptom evolution. Clinical studies report that **30–50%** of patients with DPN fail to achieve meaningful symptom relief with first-line therapies, contributing to delayed improvement, unnecessary medication exposure, and increased risk of irreversible nerve damage.

Why Existing Machine Learning Approaches Fall Short

Recent years have seen increasing interest in applying machine learning to DPN and related neurological disorders. While many models report high classification accuracy, their **clinical translation remains limited**.

Major shortcomings include:

- **Single-modality dependence**, often relying exclusively on clinical scores or EMG data
- **Black-box architectures** with limited interpretability
- Lack of **biological grounding** linking pathophysiological mechanisms to functional impairment
- **Poor generalizability** across patient populations
- Minimal alignment with **real-world clinical workflows**

Importantly, most existing models focus on **outcome prediction alone**, without addressing *why* patients differ or *how* predictions should inform treatment decisions. This disconnect has restricted clinician trust and hindered bedside adoption, despite promising numerical performance.

Unmet Clinical Need

Together, these limitations highlight a critical gap in current DPN care: the absence of a framework capable of **decoding disease heterogeneity in a mechanistically grounded, interpretable, and clinically actionable manner**. Addressing this gap is essential for improving **early prognosis**, enabling **personalized treatment strategies**, and reducing the long-term burden of complications associated with diabetic peripheral neuropathy.

2. Rationale & Significance

2.1 Why Stratification Is the Missing Layer

The central challenge in Diabetic Peripheral Neuropathy is no longer disease detection, but **explaining and managing inter-patient heterogeneity**. Large clinical cohorts demonstrate that patients with comparable diabetes duration, glycemic exposure, and symptom severity can follow **widely divergent neuropathic trajectories**, ranging from slow sensory decline to rapid motor deterioration and irreversible disability. Longitudinal studies indicate that progression rates and treatment responses may vary by **>2–3×** between clinically similar patients, underscoring the inadequacy of severity-based classification alone.

Current clinical stratification relies predominantly on **symptom intensity, nerve conduction abnormalities, or functional impairment scores**, which capture disease presence but fail to encode **why progression differs**. This limitation has direct clinical consequences. Across multiple therapeutic classes, including neuropathic pain agents and non-pharmacological interventions, approximately **30–50% of patients** fail to derive meaningful benefit from first-line treatment. The absence of biologically grounded stratification perpetuates **reactive, trial-and-error care**, delaying symptom control and increasing cumulative nerve damage.

Stratification based on underlying mechanisms—rather than surface severity—represents the critical bridge between disease biology and effective intervention. However, existing stratification efforts remain fragmented, focusing on isolated signals without integrating biological drivers and functional expression into a unified framework.

2.2 Why Mechanism-Based Stratification Remains Unresolved

Despite advances in precision medicine, DPN stratification has not translated into routine clinical practice. Key unresolved barriers include:

- **Biological signals are studied in isolation**, limiting their explanatory power
- **Physiological readouts** quantify dysfunction but do not contextualize causality
- **Existing computational models prioritize prediction accuracy**, not interpretability or mechanistic insight

As a result, no widely adopted approach currently links **pathophysiological drivers (e.g., inflammatory and oxidative stress pathways)** to **functional nerve impairment** in a manner that supports clinical reasoning and decision-making. This disconnect explains why high-performing models often fail to influence real-world treatment selection.

2.3 Scientific Rationale for an Integrative Patient Representation

To resolve DPN heterogeneity, stratification must operate at the intersection of **biological risk** and **functional manifestation**. Isolated biomarkers, whether genetic scores, cytokine levels, or EMG indices, lack sufficient context when interpreted independently. Clinicians require representations that are **integrative, compact, and interpretable**, rather than fragmented numerical outputs.

A unified latent representation enables:

- Encoding **how biological dysregulation translates into nerve dysfunction**
- Capturing **disease dynamics**, not just static severity
- Supporting **comparability across patients**, despite heterogeneous inputs

The absence of such representations in current DPN literature represents a major methodological gap. Most existing frameworks stop at outcome prediction, offering limited insight into disease structure or progression logic.

2.4 Clinical and Societal Significance

The consequences of unresolved DPN heterogeneity extend beyond individual patients. Poor stratification directly contributes to:

- **Delayed therapeutic benefit**, increasing cumulative disability
- **Escalating healthcare costs**, driven by ineffective treatment cycling
- **Higher complication rates**, including foot ulceration and amputation

From a societal perspective, even modest improvements in stratification could yield disproportionate benefit. Reducing trial-and-error treatment by **30–60%**, improving early risk identification by **>70%**, or enhancing personalization accuracy to **>95%** has the potential to substantially reduce long-term morbidity and healthcare expenditure at population scale.

2.5 Translational Significance: From Stratification to Action

Critically, stratification alone is insufficient unless it informs **actionable clinical decisions**. The lack of translation from computational outputs to intervention strategies has limited adoption of advanced models in DPN care. A stratification framework that directly supports **treatment selection, risk management, and adaptive intervention planning** addresses this final translational bottleneck.

By centering stratification as the core innovation, rather than an auxiliary analytic step, this work positions heterogeneity decoding as a practical clinical tool—transforming complex biological signals into **structured, interpretable, and decision-ready insights**. This shift represents a necessary evolution in DPN management, aligning scientific understanding with real-world clinical and societal needs.

3. Research Objectives & Expected Outcomes

3.1 Overall Aim

To **reframe the understanding of Diabetic Peripheral Neuropathy (DPN)** by uncovering **latent biological–functional patterns** that are not observable through conventional symptom-based or severity-driven assessment, enabling **mechanism-informed patient stratification** rather than surface-level disease categorization.

3.2 Specific Objectives

The specific objectives of this study are to:

- 1. **Integrate biological mechanisms and functional nerve signals** into a unified patient representation that reflects how underlying dysregulation manifests clinically.
- 2. **Identify reproducible latent stratification patterns** that explain inter-patient heterogeneity beyond conventional severity scores.
- 3. **Enable interpretable stratification axes** that support clinical reasoning rather than black-box prediction.
- 4. **Link stratification outputs to prognosis and risk profiling**, supporting earlier and more confident clinical decision-making.

3.3 Expected Clinical Outcomes

* Expected outcomes are evaluated using internal cross-validation and benchmarked against conventional assessment paradigms.

Table 1. Expected Clinical Impact

Clinical Dimension	Outcome
Treatment Trial-and-Error	↓ ≈ 60% reduction in ineffective initial treatment pathways
Risk of Major Complications	↓ ≈ 35% via earlier identification of high-risk progression patterns
Treatment Cost Burden	↓ ≈ 45% through avoidance of unnecessary escalation
Personalization Accuracy	↑ ≈ 95–98% patient-level stratification consistency
Early Risk Identification	↑ ≈ 73% improvement in detecting adverse trajectories
Treatment Response Estimation	↑ ≈ 26% improvement in individualized response prediction
Clinical Decision Consistency	↑ ≈ 92% agreement across stratification-driven decisions

4. Research Questions & Hypotheses

Rather than asking whether machine learning can *predict* diabetic peripheral neuropathy, this study is driven by a more fundamental question:

Can DPN be computationally interpreted as a heterogeneous biological process rather than a single clinical diagnosis?

Building on this premise, the research addresses the following questions:

1. **Interpretability beyond detection**
 - Can machine learning move beyond symptom detection and severity scoring to **decode the biological logic underlying DPN heterogeneity**?
2. **Hidden structure of disease variability**
 - Can a **unified, patient-specific representation** of inflammatory, oxidative, and neurofunctional signals uncover **latent disease phenotypes** that are not captured by conventional clinical classification?
3. **From signals to meaning**
 - Can complex multimodal biomarkers be transformed into a **structured, interpretable disease representation** that clinicians can reason about, rather than isolated numerical outputs?

These questions shift the role of ML from a *diagnostic classifier* to a **computational framework for disease understanding**, positioning stratification as an interpretive task rather than a predictive endpoint.

Testable Hypotheses

A) Disease Interpretation Hypothesis

Machine learning models that integrate **pathophysiological and electrophysiological biomarkers** can function as **computational interpreters of disease biology**, revealing structured heterogeneity that is not observable through symptom-based assessment alone.

B) Latent Representation Hypothesis

A learned **Personalized Vector (PV)** can encode patient-specific biological–functional interactions, enabling reproducible stratification into **mechanistically distinct neurophenotypes** beyond conventional severity categories.

C) Clinical Meaningfulness Hypothesis

Representing patient state as a **personalized disease language**, rather than isolated features, improves the **consistency, interpretability, and clinical utility** of stratification-driven decision-making.

5. Methodology

5.1 Experimental Design & Engineering Architecture

The proposed system is engineered to address a specific failure mode in current DPN modeling: **the inability of existing models to jointly encode disease mechanisms and neurofunctional expression without collapsing one modality into the other.**

Rather than treating biomarkers, EMG signals, and clinical metadata as interchangeable features, the architecture explicitly separates **cause-level signals** from **effect-level signals**, then re-integrates them through controlled representation learning.

5.1.1 Mechanism and Metadata Encoding (Autoencoder Block)

Why an Autoencoder?

Pathophysiological markers (TNF- α , MDA) and clinical metadata are **low-dimensional but noisy**, highly correlated, and sensitive to scale and distributional shifts across cohorts. Directly feeding these variables into downstream classifiers leads to unstable decision boundaries and overfitting.

An **Autoencoder (AE)** is therefore used to:

- Enforce an **information bottleneck**, preserving only biologically relevant variance
- Remove redundancy between inflammatory, oxidative, and contextual variables
- Produce a compact ***mechanism_latent*** that represents *biological load*, not raw measurements

This block encodes:

- **TNF- α** (inflammatory signaling)
Patient-level serum data obtained via direct collaboration from:
El-Badawy et al., 2017 — *Egyptian Rheumatology and Rehabilitation*
<https://doi.org/10.4103/1110-161X.205663>
El-Sheikh et al., 2019 — *The Egyptian Journal of Neurology, Psychiatry and Neurosurgery*
<https://doi.org/10.1186/s41983-019-0080-0>
- **MDA** (Malondialdehyde)
Feldman et al., 2017 — *Diabetes/Metabolism Research and Reviews*
<https://doi.org/10.1002/dmrr.287>
- **Calibrated MDA** (oxidative stress modifier; see Section 5.2.4)
- Metadata (age, sex, ethnicity, diabetes duration, DPN duration, severity grade, pain score)

Engineering rationale:

The AE ensures that downstream fusion operates on *biological state*, not heterogeneous raw inputs.

5.1.2 Neurofunctional Encoding (1D-CNN + BiLSTM Block)

Why CNN + BiLSTM instead of summary EMG features or static ML?

Electromyography reflects both:

1. **Local signal morphology** (motor unit potential shape, amplitude, integrity)
2. **Temporal recruitment behavior** (firing consistency, degradation patterns)

Single-layer CNNs capture morphology but ignore sequence context.

RNN-only models capture sequence but lose spatial signal structure.

The hybrid design combines:

- **1D-CNN** → extracts local motor unit patterns
- **BiLSTM** → models bidirectional temporal dependencies in recruitment dynamics

Input features:

- **MUNE** – motor unit loss
- **MUI** – recruitment efficiency
- **MUAPI** – action potential integrity

Patient-level EMG cohorts obtained from:

Alkotami et al., 2024

<https://doi.org/10.1186/s41983-024-00923-8>

Engineering rationale:

This block explicitly models *functional nerve expression* as a time-dependent process, not a static score.

The output is an **EMG_latent** that encodes neurofunctional degradation trajectories.

5.1.3 Cross-Modal Fusion and Disease Representation (Transformer-based cHVAE)

Why Transformer-based conditional VAE (cHVAE)?

Simple concatenation of latents fails in multimodal medical data due to:

- Modality dominance (EMG overwhelms biomarkers or vice versa)
- Loss of interaction structure between biology and function

The fusion module addresses this via:

- **Transformer attention**, allowing the model to learn *which biological factors attend to which functional signals*

- **Conditional VAE structure**, enforcing a **structured latent space** suitable for stratification rather than arbitrary embeddings

Inputs:

- mechanism_latent (biology + context)
- EMG_latent (functional impairment)

Output:

- Personalized Vector (PV) — a patient-specific disease representation encoding *how biological stress manifests as nerve dysfunction*

Engineering rationale:

The PV is not optimized for prediction alone, but for **disease state consistency across multiple downstream tasks**.

5.1.4 In-Vitro MDA Calibration (Experimental Engineering Layer)

Oxidative stress is known to influence DPN progression, yet patient-level MDA values suffer from assay variability and cohort-dependent scaling.

To prevent oxidative stress from acting as a noisy feature, an **independent in-vitro calibration layer** was implemented.

Location: National Research Centre (NRC), Egypt.

Purpose:

- Establish a **stable oxidative stress axis** decoupled from patient measurement noise

Calibration defines:

- Low / moderate / high oxidative stress bands
- Measurement stability
- Valid dynamic range for continuous modeling

Table 2. Direct system-level effects:

System Component	Effect of Elevated Calibrated MDA
Personalized Vector (PV)	Increased dispersion → higher heterogeneity
Pathophysiological Fragmentation Index (PFI)	Higher oxidative decoupling across mechanisms
Management Risk Window	Shortened effective intervention window
Outcome Stability	Wider uncertainty bounds

5.1.5 Decision Heads and System Outputs

The PV feeds two parallel decision paths:

- 1. **Multi-head MLP classifier**
Produces **three orthogonal stratification axes**:
 - Axis 1: progression dynamics
 - Axis 2: dominant mechanism
 - Axis 3: symptom–damage alignment
- 2. **Multi-task MLP regressor**
Estimates:
 - Nerve biological age
 - Predicted outcome trajectory

Additional derived constructs:

- **Management Risk Window**
- **Pathophysiological Fragmentation Index (PFI)**
- **Mechanisms–to–Muscle Interaction Map** (via SHAP + GNN attribution)

Engineering rationale:
All outputs originate from the *same representation*, forcing internal coherence between stratification, prognosis, and interpretability.

5.2 System Outputs & Why They Matter (Engineering → Clinical)

Table 3. Outputs and Clinical Relevance:

Output	Engineering Role	Clinical Value
Personalized Vector	Unified mechanism function embedding	Patient-specific disease signature
Axis Stratification	Structured heterogeneity decomposition	Progression · Mechanism Symptom–damage
Management Risk Window	PV-conditioned decision interval	Optimal timing for intervention
Interaction Map	Mechanism→function attribution graph	Interpretable cause–effect links
Nerve Biological Age	PV-anchored regression	Functional aging beyond chronology
Predicted Outcome	Uncertainty-aware prognosis	Trajectory-guided care planning
Pathophysiological Fragmentation Index	Mechanism–function decoupling score	Non-uniform damage severity
VR-Guided Intervention	Closed-loop PV-driven module	Personalized neuromodulation (Phase II)

6. Human Participants

6.1 Participant Characteristics

This study utilized **retrospective human clinical data** derived from previously conducted research studies on diabetic peripheral neuropathy. The datasets included adult participants diagnosed with diabetes and clinically confirmed DPN, representing a spectrum of disease duration, symptom severity, and neurofunctional impairment.

Participant-level data were accessed in a **coded (non-public) form**, allowing clinical variables and biomarker measurements to be analyzed while maintaining participant confidentiality.

6.2 Recruitment and Informed Consent

No new participants were recruited for this research. All individuals included in the source datasets had previously provided **written informed consent** permitting the use of their clinical, biological, and electrophysiological data for **secondary research purposes**.

Data access was granted through direct collaboration with the original study investigators, in accordance with the consent terms approved under the original study protocols. The current work constitutes **secondary analysis within the scope of the original participant consent**, and therefore did not require additional participant contact or re-consenting.

6.3 Privacy, Data Handling, and Access Control

Although the data were not fully anonymous at the point of original collection, all datasets were **coded and stripped of direct personal identifiers** prior to analysis in this study.

Data protection measures included:

- Use of **coded identifiers** instead of personal information
- Secure, access-restricted data storage
- Analysis limited strictly to approved research objectives
- No data was shared beyond the authorized research scope.

Ethical and Regulatory Status

This research represents a **secondary analysis of previously consented human data**, conducted without direct participant interaction or intervention. As data use was explicitly permitted under the original informed consent agreements and no new data collection occurred, **additional institutional ethical approval was not required** for this phase of the study.

7. Data Analysis

This section defines the **quantitative and analytical procedures** used to evaluate model behavior, stability, and clinical reliability across stratification and prognosis tasks.

7.1 Statistical Characterization of Inputs

Continuous variables were analyzed using:

- Mean $\pm$ SD
- Interquartile range (IQR)
- Coefficient of variation (CV%)
- ANOVA test

Observed variability:

- TNF- α : CV $\approx$ **22–28%**
- MDA: CV $\approx$ **18–24%**
- EMG features: CV $\approx$ **25–35%**, reflecting heterogeneous neurofunctional degradation

Correlation screening (Pearson $|r|$):

- TNF- α $\leftrightarrow$ EMG impairment: **$r \approx 0.48\text{--}0.55$**
- MDA $\leftrightarrow$ functional dispersion: **$r \approx 0.51\text{--}0.60$**

ANOVA formulation :

$$F = \frac{\text{between-group variance}}{\text{within-group variance}}$$

Hypotheses

$$H_0 : \mu_1 = \mu_2 = \dots = \mu_k$$

$$H_a : \exists i \neq j \text{ such that } \mu_i \neq \mu_j$$

Table 4. Observed Statistical Separation

Variable	F-statistic (range)	p-value
PV dispersion	18.2 – 26.7	$p < 0.001$
PFI	21.4 – 34.1	$p < 0.001$
Nerve Biological Age	16.9 – 22.3	$p < 0.001$
Outcome score	14.6 – 19.8	$p < 0.001$

These results reject the null hypothesis, indicating significant separation between stratified groups. informed representation learning but did not drive feature elimination.

7.2 Stratification Performance Analysis (Classification)

Stratification outputs were evaluated per axis and jointly using macro-averaged metrics.

Primary Metrics

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad \text{Macro-F1} = \frac{1}{K} \sum_{k=1}^K \frac{2 \cdot P_k \cdot R_k}{P_k + R_k} \quad \text{ROC-AUC}_{OVR} = \frac{1}{K} \sum_{k=1}^K AUC_k$$

Table 5. Observed ranges (test set):

Metric	Value Range
Accuracy	0.96 – 0.98
Macro-F1	0.97 – 0.99
ROC-AUC (OVR)	0.97 – 1.00
Cohen's κ	0.95 – 0.98

Axis-wise performance deviation < **2.5%**, indicating balanced stratification across progression, mechanism, and symptom–damage axes.

7.3 Prognosis Performance Analysis (Regression)

Regression heads were evaluated using absolute error, variance explanation, and correlation strength.

Error Metrics

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i| \quad RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}$$

Association Metrics

$$R^2 = 1 - \frac{\sum (y - \hat{y})^2}{\sum (y - \bar{y})^2} \quad \rho_{Spearman} = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

Table 6. Observed performance:

Metric	Value
MAE	0.037
RMSE	0.061
R ²	≈ 0.95
Pearson r	≈ 0.96
Spearman ρ	≈ 0.98

Residual dispersion remained symmetric around zero with no heteroscedastic drift.

7.4 Calibration Analysis

Calibration quality was quantified using:

$$ECE = \sum_{m=1}^M \frac{|B_m|}{N} |acc(B_m) - conf(B_m)| \quad MCE = \max_m |acc(B_m) - conf(B_m)|$$

Table 7. Observed calibration stability:

Metric	Range
ECE	0.002 – 0.006
MCE	0.15 – 0.24

Reliability curves demonstrated near-diagonal alignment, confirming confidence–accuracy consistency.

7.5 Robustness and Stress Testing

Robustness was evaluated under controlled perturbations.

Noise Injection

Gaussian noise ($\sigma \in [0, 5]$) applied to EMG inputs:

- $\Delta\text{Macro-F1} \leq 3\%$ at $\sigma = 5$

Missingness Stress

Progressive masking up to **40%** of features:

- Performance degradation $< 2\%$
- No axis collapse observed

Boundary Sensitivity

Error density concentrated near decision margins ($p \approx 0.55\text{--}0.65$), validating uncertainty-aware outputs.

7.6 Uncertainty Quantification

Bootstrap resampling (1,000 iterations) was used to estimate confidence intervals:

$$95\% \text{ CI} = \hat{\mu} \pm 1.96 \cdot \frac{\sigma}{\sqrt{N}}$$

All major metrics demonstrated **stable 95% confidence** bounds, supporting reproducibility.

7.7 Interpretability and Error Decomposition

- SHAP attribution quantified feature influence magnitude
- GNN interaction maps decomposed mechanism $\rightarrow$ muscle pathways
- Entropy–margin analysis identified low-confidence, high-risk predictions

These analyses ensured that **errors were structured, explainable, and clinically interpretable**.

8. Risk Assessment & Safety Considerations

This section outlines the **actual procedural and technical risks** associated with the study components, along with the measures implemented to minimize them.

8.1 Biological Sample Handling (Serum Analysis)

Serum-based biomarker analysis (TNF- α and MDA) involved **standard venous blood collection** performed as part of the original clinical studies.

Potential risks:

- Mild discomfort or bruising at the puncture site
- Minimal risk of infection associated with routine blood draw

Risk mitigation:

- All blood sampling was conducted by qualified medical personnel
- Use of sterile, single-use collection materials
- Adherence to established clinical laboratory safety procedures

No additional blood samples were collected specifically for this study.

8.2 Electrophysiological Assessment (EMG)

Electromyography assessments were performed as part of routine neurophysiological evaluation in the source studies.

Potential risks:

- Temporary discomfort during electrode placement
- Minor muscle soreness following examination

Risk mitigation:

- EMG procedures conducted by trained neurophysiology specialists
- Use of clinically approved EMG equipment
- Application of established neuromuscular testing protocols

No experimental or invasive EMG techniques were introduced.

8.3 In-Vitro Oxidative Stress Calibration

The in-vitro calibration of oxidative stress (MDA) was performed independently of patient samples.

Safety considerations:

- Conducted under **standard laboratory safety protocols**
- No human-derived biological material was used
- Procedures were carried out at the **National Research Centre (NRC), Egypt**, in accordance with institutional laboratory regulations

This component posed **no risk to human participants**.

8.4 Digital Application and Data Handling (Future Deployment Phase)

A future mobile application layer is proposed to support personalized clinical feedback and intervention guidance.

Potential risks:

- Unauthorized access to patient-level outputs
- Data misuse during transmission or storage

Planned safeguards:

- Use of **pseudonymized user identifiers**
- Separation of clinical outputs from personal identity data
- Encrypted data transmission and local secure storage
- No storage of raw biological signals on the device
- Access restricted to authorized clinical or research use only

The application is designed to function as a **decision-support interface**, not a diagnostic replacement.

Overall Risk Profile

All identified risks are **minimal, well-characterized, and routinely managed** within standard clinical and laboratory practice. No experimental procedures, unvalidated devices, or high-risk interventions were introduced at any stage of the study.

9. Ethical Considerations

This study was conducted in accordance with established principles of **responsible research and ethical data use**. All human-related data analyzed in this work originated from prior clinical studies in which participants provided informed consent for research use of their data.

The current research involved **secondary analysis only**, with no direct participant interaction, no intervention, and no alteration to standard clinical care.

The study design prioritizes **interpretability, transparency, and clinical responsibility**, ensuring that computational outputs are intended to support—not replace—clinical judgment.

Future translational components, including digital decision-support interfaces, are designed to operate under clinician oversight and comply with applicable data protection and ethical standards.

10. Translational Consideration

This work is designed as a **translationally oriented computational framework**, with outputs structured to align with real-world clinical decision workflows rather than experimental endpoints.

The proposed system operates as a **clinical decision-support layer**, integrating routinely measurable biomarkers and electrophysiological data without altering existing diagnostic or treatment pathways. All outputs—including stratification axes, risk windows, and prognostic indicators—are intended to **support clinician judgment**, not to replace it.

From an implementation perspective:

- The system relies on **standard clinical inputs** (serum biomarkers, EMG, metadata), enabling integration into existing care settings.
- Computational outputs are modular and interpretable, facilitating **incremental adoption** rather than full system replacement.
- Risk-sensitive outputs (e.g., management risk window) are designed to inform **timing and prioritization**, not autonomous intervention.

Future extensions, including personalized digital interfaces and VR-guided neuromodulation, are explicitly positioned as **adjunctive tools** subject to clinical oversight and ethical governance. No direct therapeutic action is automated at any stage.

Overall, the framework prioritizes **clinical feasibility, interpretability, and safety**, supporting gradual translation from computational modeling to practical decision support in DPN management.

11 .Bibliography

Alkotami, A., et al. (2024). Electrophysiological evaluation of diabetic peripheral neuropathy. *The Egyptian Journal of Neurology, Psychiatry and Neurosurgery*, 60, 44.

<https://doi.org/10.1186/s41983-024-00923-8>

Pop-Busui, R., Boulton, A. J. M., Feldman, E. L., et al. (2017). Diabetic neuropathy: A position statement by the American Diabetes Association. *Diabetes Care*, 40(1), 136–154.

<https://doi.org/10.2337/dc16-2042>

Feldman, E. L., Nave, K.-A., Jensen, T. S., & Bennett, D. L. H. (2017). New horizons in diabetic neuropathy. *Nature Reviews Neurology*, 13, 541–553.

<https://doi.org/10.1038/nrneurol.2017.126>

Zhang, Y., Xu, L., Zhu, Q., et al. (2023). Diabetic peripheral neuropathy: Pathogenetic mechanisms and healthcare burden. *Frontiers in Endocrinology*, 14, 1265372.

<https://doi.org/10.3389/fendo.2023.1265372>

Guo, Q.-F., He, L., et al. (2022). Virtual reality for neurorehabilitation: A bibliometric analysis. *Frontiers in Public Health*, 10, 1042618.

<https://www.frontiersin.org/articles/10.3389/fpubh.2022.1042618/full>

Chew, S. M., Dua, S., & Venkataraman, K. (2025). Predictors of incident diabetic peripheral neuropathy: A systematic review. *Reviews in Endocrine and Metabolic Disorders*, 26, 659–677.

<https://doi.org/10.1007/s11154-025-09973-6>

Zhong-Chong, Z., & Souayah, N. (2025). Crumbling pathogenesis and biomarkers for diabetic peripheral neuropathy. *Biomedicines*, 13(2), 413.

<https://doi.org/10.3390/biomedicines13020413>

El-Badawy, E. A., et al. (2017). Serum TNF- α levels in diabetic peripheral neuropathy. *Egyptian Rheumatology and Rehabilitation*, 44, 27–33.

<https://doi.org/10.4103/1110-161X.205663>

El-Sheikh, A. A., et al. (2019). Inflammatory biomarkers in diabetic peripheral neuropathy. *The Egyptian Journal of Neurology, Psychiatry and Neurosurgery*, 55, 43.

<https://doi.org/10.1186/s41983-019-0080-0>

Kiziltunc, A., et al. (2014). Serum prolidase enzyme activity and oxidative stress levels in patients with diabetic neuropathy. *Endocrine*, 47, 146–151.

<https://doi.org/10.1007/s12020-013-0136-3>

Feldman, E. L., et al. (2017). Oxidative stress and diabetic neuropathy: Pathophysiological mechanisms and treatment perspectives. *Diabetes/Metabolism Research and Reviews*, 33(5), e287.

<https://doi.org/10.1002/dmrr.287>

Sun, C.-F., Lin, Y.-H., Ling, G.-X., et al. (2025). Systematic review and critical appraisal of predictive models for diabetic peripheral neuropathy. *World Journal of Diabetes*, 16(4), 101310.

<https://doi.org/10.4239/wjd.v16.i4.101310>

Zhao, Y., Chen, C., et al. (2023). Explainable artificial intelligence in clinical risk prediction models. *Journal of Clinical Medicine*, 12(6), 2189.

<https://doi.org/10.3390/jcm12062189>

Guo, Q.-F., He, L., et al. (2022). Virtual reality for neurorehabilitation: A bibliometric analysis. *Frontiers in Public Health*, 10, 1042618.

<https://doi.org/10.3389/fpubh.2022.1042618>

Chenais, N., & Görgen, A. (2024). Immersive interfaces for clinical applications. *Frontiers in Neurobotics*, 18, 1362444.

<https://doi.org/10.3389/fnbot.2024.1362444>

International Diabetes Federation. (2025). *IDF Diabetes Atlas* (11th ed.).

<https://diabetesatlas.org>
